
SafeDrug: A Benchmark Dataset for Safety-Critical Pharmacological Reasoning in LLMs

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Large language models (LLMs) have shown strong performance on biomedical
2 tasks, yet evaluating their reasoning in safety-critical pharmacological contexts
3 requires large-scale, structured datasets. We present SafeDrug, a benchmark dataset
4 for systematic evaluation of pharmacological reasoning across multiple task di-
5 mensions. SafeDrug integrates heterogeneous sources, including adverse drug
6 event records, drug–drug interaction knowledge, literature-derived evidence, and
7 population-specific information such as child, adult, and older adult groups, into
8 a unified reasoning-oriented QA format. The benchmark comprises two com-
9plementary subsets. **SafeDrug-large** provides broad coverage at scale, while
10 **SafeDrug-small** is a high-quality human-annotated subset for precise evalua-
11tion. It supports tasks such as prediction, risk assessment, mechanism expla-
12nation, evidence-grounded QA, and drug replacement across both single-drug
13and multi-drug settings. Our evaluation shows that LLMs capture outcome-
14level associations but struggle with mechanistic reasoning, evidence alignment,
15and population-aware predictions. SafeDrug provides a foundation for repro-
16ducible benchmarking and advances research on evidence-grounded and popula-
17tion-aware pharmacological reasoning in LLMs. The code and data are available at:
18 <https://anonymous.4open.science/r/SafeDrug-0544>

19 1 Introduction

20 Large language models (LLMs) have recently demonstrated strong performance across a wide
21 range of biomedical and clinical tasks [1, 2], including medical question answering [3, 4], literature
22 understanding [5], and knowledge retrieval [6]. These capabilities suggest the potential of LLMs
23 to support real-world healthcare applications. However, many of these tasks primarily evaluate
24 factual recall or short-form reasoning, leaving open the question of whether LLMs can reliably
25 perform *safety-critical pharmacological reasoning*, such as assessing adverse drug events (ADEs [7]),
26 identifying drug–drug interactions (DDIs [8, 9]), and providing mechanism-grounded explanations.

27 Pharmacological safety assessment is fundamentally more challenging than standard biomedical
28 QA [1], which requires multi-step reasoning over heterogeneous knowledge sources [10, 11, 12],
29 including drug properties, biological mechanisms, and clinical evidence. In addition, real-world
30 scenarios often involve *compositional complexity*, such as multi-drug interactions, and *population*
31 *heterogeneity*, where risk profiles differ across patient groups (e.g., children, adults, and older adults).
32 Errors in such settings may lead to incorrect or unsafe clinical recommendations, making reliability
33 and interpretability essential. Despite growing interest in evaluating LLMs in the biomedical domain,
34 existing benchmarks remain insufficient for assessing these capabilities [13, 7, 14]. Most prior
35 datasets focus on single-step question answering or classification tasks [7, 15] and lack key aspects
36 required for safety-critical reasoning [13], such as mechanism-level supervision [16], explicit evidence

Benchmark	Domain	Reasoning	Long Chain	Mechanistic	Population	Evidence	Therapy
MedQA [20]	General Medicine	1-step QA	✗	✗	✗	✗	✗
PubMedQA [21]	Biomedical	1-step QA	✗	✗	✗	✓	✗
BioASQ [22]	Biomedical	1-step QA	✗	✗	✗	✓	✗
DrugBank-QA [10]	Discovery	1-step QA	✗	✓	✗	✗	✗
DDI-Ben [13]	Interaction	Classification	✗	✗	✗	✗	✗
SafeDrug	Drug safety	≥2-step QA	✓	✓	✓	✓	✓

Table 1: Comparison of SafeDrug with existing biomedical and drug safety benchmarks. SafeDrug differs from prior benchmarks by jointly supporting multi-step pharmacological reasoning, mechanism-level evaluation, population-aware assessment, evidence alignment, and therapy-level decision tasks. ✓ indicates that the capability is supported, while ✗ indicates that it is not supported.

grounding, and population-aware evaluation. As a result, it remains unclear whether current LLMs genuinely understand pharmacological mechanisms or merely rely on superficial correlations in data.

To address this gap, we introduce **SafeDrug**, a comprehensive benchmark for evaluating pharmacological safety reasoning in large language models. SafeDrug organizes tasks into a hierarchical structure spanning prediction and reasoning, including ADE prediction [17], DDI detection [18, 19], risk assessment, mechanism explanation, evidence-grounded question answering, and drug replacement. The benchmark integrates heterogeneous pharmacological datasets from multiple sources and unifies them into a reasoning-oriented QA format. Importantly, SafeDrug explicitly incorporates *population-specific evaluation* and supports *multi-drug scenarios*, enabling a more realistic assessment of model behavior in complex clinical settings.

Using SafeDrug, we conduct a systematic evaluation of state-of-the-art LLMs and uncover several key limitations. First, while models achieve competitive performance on prediction tasks, they struggle with mechanism explanation and evidence-grounded reasoning, revealing a gap between knowledge memorization and mechanistic understanding. Second, performance degrades significantly with increasing reasoning complexity, particularly in multi-drug settings, indicating limited compositional reasoning ability. Third, models often fail to faithfully utilize supporting evidence, producing hallucinated or weakly grounded justifications. Finally, we observe substantial performance disparities across patient populations, suggesting that current models lack robustness in underrepresented groups.

Our work highlights fundamental challenges in applying LLMs to safety-critical pharmacological tasks and provides a benchmark for systematically evaluating and understanding these limitations. We summarize our main contributions as follows: (1) We introduce **SafeDrug**, a comprehensive benchmark dataset for evaluating safety-critical pharmacological reasoning in large language models, consisting of a large-scale collection (SafeDrug-large) and a high-quality human-annotated subset (SafeDrug-small). (2) We design a hierarchical task framework spanning prediction, risk assessment, mechanistic reasoning, evidence grounding, and decision-making under multi-drug and population-specific settings. (3) We conduct a large-scale evaluation of state-of-the-art LLMs, revealing systematic gaps between outcome-level prediction and deeper reasoning, particularly in mechanism understanding, evidence alignment, and population robustness. (4) We provide detailed diagnostic analyses, including capability breakdown and error taxonomy, to support future research on reliable, interpretable, and population-aware pharmacological reasoning.

2 Benchmarking Adverse Drug Events with SafeDrug

We introduce **SafeDrug**, a comprehensive benchmark for evaluating safety-critical pharmacological reasoning in LLMs. As illustrated in Figure 1, SafeDrug is designed to move beyond conventional prediction-based evaluation by incorporating structured reasoning, evidence grounding, and population-aware decision-making into a unified framework.

2.1 Benchmark Tasks

SafeDrug organizes pharmacological tasks into a hierarchical structure that spans *classification* and *reasoning* (Figure 1a). This design reflects the increasing complexity of drug safety assessment, from outcome prediction to structured, multi-step reasoning.




Figure 1: SafeDrug is a benchmark for evaluating safety-critical pharmacological reasoning in large language models. It covers a hierarchy of tasks from prediction to multi-step reasoning (a), integrates heterogeneous ADE/DDI data with population awareness into a unified QA format (d), and evaluates diverse LLMs using multiple metrics for both accuracy and reasoning quality (b–c). The pipeline standardizes data, performs model inference, and conducts multi-dimensional evaluation (e).

Classification Tasks. At the base level, SafeDrug includes outcome-oriented prediction tasks: (i) *prediction* and (ii) *risk assessment*. The prediction task focuses on identifying potential adverse drug events (ADEs) or interactions, while risk assessment requires estimating the severity or likelihood of such events. These tasks primarily evaluate whether models can capture statistical associations between drugs and clinical outcomes. However, they do not require an understanding of underlying mechanisms or justification of predictions.

Reasoning Tasks. To evaluate deeper pharmacological understanding, SafeDrug introduces reasoning-centric dimensions such as *evidence grounding*, which requires models to support predictions with relevant evidence, ensuring both correctness and justification. This is critical in safety-sensitive settings, where unsupported outputs may lead to unreliable conclusions. To assess this capability, we design two complementary tasks: evidence selection and evidence generation. (iv) *mechanistic reasoning* evaluates whether models can explain drug effects and interactions based on biological or pharmacological mechanisms (e.g., metabolism or toxicity pathways). This dimension goes beyond correlation to assess mechanistic understanding. To this end, we include both mechanistic classification and generation tasks. (v) *compositional reasoning* focuses on multi-drug scenarios, where interactions emerge from combinations of drugs. These tasks require models to reason over multiple entities and their interactions for prediction and risk assessment, reflecting real-world settings such as polypharmacy. Together, these dimensions define a progression from surface-level prediction to structured reasoning. In particular, SafeDrug explicitly models the transition from *single-drug* to *multi-drug* settings, increasing reasoning complexity and enabling systematic evaluation of model capabilities under realistic pharmacological conditions.

2.2 Data Integration

SafeDrug integrates heterogeneous pharmacological data sources into a unified reasoning-oriented format (Figure 1b). More details about preprocessing and QA templates in Appendix A.3 and A.4.

ADE Data for Single Drug. We incorporate large-scale adverse drug event (ADE) datasets, including SIDER [23], OFFSIDES [24], CT-ADE [7], and pediatric-specific resources [25]. These datasets

provide drug–adverse effect associations, frequency statistics, and clinical reports, capturing relationships between drug exposure and observed outcomes. ADE data primarily support *prediction* and *risk assessment* tasks, offering a statistical foundation for identifying potential adverse events. However, such associations do not encode causal mechanisms or supporting evidence, and thus mainly assess the knowledge recall capability of LLMs rather than their reasoning ability.

ADE Data for multiple Drugs. We incorporate curated drug–drug interaction (DDI) datasets, including DrugBank [10], HODDI [26], and DDInter2.0 [15]. These resources provide structured information on interaction types, severity levels, and, in some cases, underlying pharmacological mechanisms (e.g., metabolic pathways or transporter effects). In particular, HODDI enables evaluation of *compositional reasoning* in multi-drug settings and offers partial support for *mechanistic reasoning*. Importantly, these datasets reflect the combinatorial nature of real-world drug usage, where interactions arise from multiple agents rather than isolated drugs, and provide sufficient coverage for evaluating LLMs on both prediction and compositional reasoning tasks.

Population-Aware Data. To capture clinical heterogeneity, SafeDrug incorporates population-specific information, including pediatric, adult, and older adult groups, from pediatric-focused resources [25] and CT-ADE [7]. These data reflect age-related differences in drug response, such as variations in metabolism, organ function, and comorbidity profiles. Population-aware data enable evaluation of whether models can adapt predictions and reasoning across patient groups, supporting more realistic and safety-critical assessment scenarios.

Evidence and Literature Data. In addition to structured datasets, SafeDrug integrates large-scale literature-derived evidence from DDInter2.0, comprising approximately 33K curated biomedical documents. We further process these sources using DeepSeek-V3.2 [27] to extract and summarize adverse event mechanisms and supporting evidence, which are then organized into reasoning-oriented QA instances. To ensure data quality, we perform stratified sampling based on the number of supporting documents per event and conduct manual verification on a subset of instances (refer to Appendix A.4). This pipeline provides reliable and structured evidence for evaluating *evidence grounding* and *mechanistic reasoning*, and supports future research on LLM fine-tuning for pharmacological reasoning.

Unified QA-based Format. All data sources are unified into a *reasoning-oriented QA format*, where each instance includes a question, structured context, and target outputs (e.g., answer, mechanism, evidence). This design enables consistent evaluation across heterogeneous tasks and supports assessment of both prediction accuracy and reasoning quality. By integrating complementary data modalities, SafeDrug facilitates comprehensive evaluation of pharmacological reasoning under realistic clinical conditions. Details of the construction process and task-specific templates are provided in Appendix A.

2.3 Benchmark Construction

The construction pipeline of SafeDrug is illustrated in Figure 1c, where heterogeneous pharmacological data are transformed into a unified and comparable evaluation framework for LLMs. The overall data distribution is shown in Figure 2, with detailed statistics provided in Appendix A.

Data Standardization and QA Construction. Raw data from heterogeneous sources (e.g., structured databases and textual evidence) are first standardized via preprocessing, including entity normalization and terminology alignment (e.g., mapping drugs and adverse events to unified vocabularies). This ensures consistency across datasets and avoids spurious discrepancies caused by format or naming differences. We then convert standardized data into a *reasoning-oriented QA format* using task-specific prompt generation, where each instance is framed as a question requiring prediction, reasoning, or decision-making with structured context. This design enables a unified evaluation interface across tasks and aligns evaluation with the interactive reasoning capabilities of modern LLMs. Each instance follows a consistent schema (e.g., `task_name`, `prompt`, `question`, `answer`), supporting parallel evaluation and fine-grained analysis of prediction accuracy, reasoning quality, and evidence alignment, while allowing extension to additional outputs such as mechanisms and supporting evidence. Detailed descriptions of mechanism labels and evidence supports are provided in Appendix A.1. Details of drug replacement (or substitution) labels are provided in Appendix A.2.

Scale and Coverage. SafeDrug consists of two complementary subsets. **SafeDrug-large** is a large-scale collection built from **7 datasets**, containing approximately **1.9M instances** across **5 task**

156 **dimensions.** **SafeDrug-small** is a high-quality human-curated subset obtained via stratified sampling
 157 of **3,000** instances based on supporting evidence counts, followed by expert verification, resulting
 158 in **2,554** high-confidence instances. In addition, SafeDrug integrates over **33K literature-derived**
 159 **evidence entries** to support evaluation of evidence grounding and mechanistic reasoning, enabling
 160 both large-scale coverage and reliable assessment of pharmacological reasoning. We provide a
 161 detailed annotation protocol for human labeling in Appendix A.4.

162 2.4 Evaluation Protocol

163 SafeDrug adopts a
 164 multi-dimensional
 165 evaluation protocol to
 166 assess both answer cor-
 167 rectness and reasoning
 168 reliability (Figure 1e).
 169 This is necessary due
 170 to the heterogeneous
 171 nature of drug safety
 172 tasks, which range
 173 from binary decisions
 174 to set prediction and
 175 free-form mechanistic
 176 explanations. Ac-
 177 cordingly, a single
 178 metric is insufficient
 179 to capture the full
 180 spectrum of ADE and
 181 DDI reasoning ability.

182 Evaluation Metrics.

183 We adopt three com-
 184plementary metric
 185families to evaluate different task types. For prediction-oriented tasks (e.g., ADE prediction,
 186DDI detection, and risk assessment), we use classification metrics such as Accuracy, Binary-F1,
 187Macro/Micro-F1, and ROC-AUC, where Macro-F1 is particularly important due to class imbalance
 188in adverse events. For tasks with multiple valid outputs (e.g., ADE sets, evidence selection, and
 189multi-drug interactions), we use set-based metrics including Set-F1, Exact Match, and Partial
 190Match to capture both correctness and coverage. For open-ended reasoning tasks (e.g., mechanism
 191explanation and evidence-grounded QA), we use generation metrics such as ROUGE, BLEU, and
 192LLM-as-a-judge to assess semantic consistency, coherence, and evidence alignment. Detailed metric
 193formulations are provided in Appendix B.1.

194 2.5 Model Coverage

195 SafeDrug evaluates a diverse set of large language models, including both API-based (closed-source
 196LLMs) and open-weight models (Figure 1d). This design is intended to assess whether safety-critical
 197pharmacological reasoning is a general capability of current LLMs or primarily emerges in the
 198strongest proprietary systems.

199 **API-based Models.** We include state-of-the-art API-based models, which typically represent the
 200strongest available general-purpose and reasoning-oriented LLMs. Evaluating these models provides
 201an estimate of the current upper bound of LLM performance on drug safety reasoning tasks. In
 202particular, they allow us to examine whether high-capacity models can reliably perform mechanism
 203explanation, evidence-grounded reasoning, and population-specific ADE assessment.

204 **Open-weight Models.** We also evaluate open-weight models to ensure reproducibility and to
 205understand how accessible models perform under the same benchmark protocol. These models are
 206important for analyzing whether pharmacological reasoning failures are shared across model families,
 207and whether smaller or openly available systems exhibit larger gaps in evidence use, mechanism
 208reasoning, and multi-drug compositionality.




Figure 2: Data statistics of SafeDrug. (a) Task distribution across prediction, risk assessment, and reasoning dimensions. (b) Population-aware data distribution over child, adult, and older adult groups. (c) Distribution of multi-drug interactions with varying numbers of drugs, reflecting polypharmacy complexity.

Table 2: Performance of representative LLMs on SafeDrug single-drug and pairwise DDI tasks. Exact Match scores show that models perform better on prediction tasks than on reasoning tasks, highlighting limitations in mechanistic understanding.

Model	ADE		Drug Interaction		Reasoning	
	Prediction	Risk	Prediction	Risk	Drug Replacement	Mechanism Prediction
Qwen2.5:7B	15.73	29.53	28.16	42.55	68.42	23.28
Qwen2.5:3B	13.90	12.39	21.42	29.31	60.34	42.49
LLaMA3.1:8B	3.06	32.94	28.96	47.15	66.40	4.22
LLaMA.2:3B	14.78	12.85	29.82	40.36	59.40	5.02
DeepSeek-V3	23.15	29.77	26.96	35.12	59.90	29.23
GPT-4o-mini	13.49	25.49	29.29	47.43	68.06	31.04
Gemini-2.5-flash	10.90	38.10	63.55	35.60	78.85	21.70
Qwen3-235B	11.05	25.10	32.55	29.00	75.70	27.00
GPT-5-mini	8.25	27.00	57.60	48.90	78.80	24.60

Table 3: Performance of representative LLMs on SafeDrug polypharmacy tasks.

Model	Multiple Drug Prediction				Multiple Drug Risk			
	3 Drugs	4 Drugs	5 Drugs	10 Drugs	3 Drugs	4 Drugs	5 Drugs	10 Drugs
Qwen2.5:7B	47.11	48.60	53.13	65.29	42.83	38.66	39.27	45.24
Qwen2.5:3B	57.60	63.30	66.42	69.42	68.33	74.71	84.83	95.24
LLaMA3.1:8B	41.25	45.05	52.30	43.80	37.39	32.25	28.20	21.43
LLaMA3.2:3b	40.32	51.94	61.64	65.29	58.93	58.42	62.45	66.67
DeepSeek-V3	60.34	64.59	65.76	66.11	63.98	69.96	76.04	91.66
GPT-4o-mini	43.74	48.83	52.08	51.23	46.21	45.57	46.42	57.14
Gemini-2.5-flash	62.60	65.05	65.90	66.11	65.20	71.60	80.20	94.04
Qwen3-235B	56.75	59.80	62.60	64.46	61.20	67.45	70.45	89.28
GPT-5-mini	58.65	63.45	65.05	66.11	50.35	53.05	56.65	69.04

3 Experiment

Additional implementation details and experimental results on SafeDrug are provided in Appendix B.

3.1 Experimental Setup

Datasets. We aggregate samples from multiple sources to construct a unified SafeDrug collection with 1.9M instances, split into train, validation, and test sets at an 8:1:1 ratio. Additionally, we create a high-precision subset of 2,554 instances (train/val/test = 1,817/331/406) for focused benchmarking. All experiments are conducted on held-out test splits to evaluate LLM performance across prediction, risk assessment, reasoning, and evidence-grounded tasks.

Models. We benchmark a representative set of nine large language models, including both open-weight and proprietary systems. Open-weight models include Qwen2.5-7B [28], Qwen2.5-3B [28], LLaMA-3.1-8B [29], and LLaMA-3.2-3B [29]. Proprietary and frontier models include DeepSeek-V3 [30], GPT-4o-mini [31], Gemini-2.5-Flash [32], Qwen3-235B-A22B [33], and GPT-5-mini [34].

Implementation Details. All experiments were conducted in a **zero-shot** setting with deterministic decoding ($T = 0$) to isolate model reasoning capabilities without task-specific fine-tuning. Open-source models were evaluated on a server with a single NVIDIA RTX 3080Ti using the Ollama framework, while closed-source models were accessed via proxy APIs (ChatAnyWhere¹). For multiple-choice tasks, answer options were shuffled using a fixed random seed (42) to mitigate positional biases. Model outputs were post-processed using pattern matching to extract final predictions and structured fields such as predicted labels, supporting evidence, and mechanistic annotations when applicable. This setup enables consistent comparison across models and tasks while controlling for variability introduced by decoding stochasticity or formatting differences. Full details of preprocessing, prompt design, and post-processing are provided in Appendix A.

¹<https://api.chatanywhere.tech>

3.2 Overall Performance on SafeDrug

We first evaluate representative LLMs on single-drug and pairwise DDI tasks using Exact Match (EM) scores, as summarized in Table 3. The evaluation covers three dimensions: *ADE prediction and risk assessment*, *DDI prediction and risk assessment*, and reasoning tasks including *drug replacement* and *mechanism prediction*. From the reported results, we have three main observations: (1) Across models, we observe that LLMs generally achieve higher performance on outcome-level prediction tasks, such as ADE and DDI prediction, compared with reasoning-oriented tasks. For instance, models like GPT-5-mini and Gemini-2.5-Flash reach EM scores above 60% on DDI prediction, while mechanism prediction remains below 30% for most models. This indicates that current LLMs are capable of capturing statistical associations between drugs and adverse events, but struggle to provide mechanistic explanations that require deeper pharmacological understanding. (2) Similarly, risk assessment tasks present a substantial challenge. Many models, including Qwen2.5 variants and LLaMA3 models, exhibit a marked drop in EM scores when predicting severity or risk labels compared with straightforward ADE prediction. This suggests that LLMs may rely on surface-level correlations rather than accurately reasoning about clinical severity. (3) Reasoning tasks such as drug replacement and mechanism prediction further reveal the limitations of current models. Even high-capacity models like GPT-5-mini achieve modest performance on mechanism prediction (24.6%) despite strong drug replacement scores (78.8%), highlighting a gap between prediction capability and mechanistic reasoning. These results underscore that SafeDrug effectively probes both predictive and reasoning capabilities, revealing systematic weaknesses in mechanistic and evidence-grounded understanding in current LLMs. Overall, the findings demonstrate that while LLMs can perform reasonably on single-step prediction tasks, they remain limited on tasks requiring multi-step reasoning, evidence alignment, and mechanistic understanding, motivating further research on improving safety-critical pharmacological reasoning.

3.3 Capability Breakdown: Prediction vs. Reasoning

To gain a more granular understanding of LLM strengths and limitations on SafeDrug, we evaluate model performance across five key capability dimensions: prediction, risk assessment, mechanistic reasoning, evidence grounding, and compositional reasoning. For each model, we compute task-specific metrics (e.g., Exact Match, Set-F1, LLM-as-a-judge scores) and aggregate them by capability to visualize relative performance, as shown in the radar plot in Figure 3. The analysis reveals systematic differences between outcome-level prediction and reasoning-oriented capabilities. Most models achieve their highest scores on prediction and compositional reasoning, indicating that LLMs can capture statistical patterns in drug-ADE associations and multi-drug correlations. By contrast, mechanistic reasoning and evidence grounding consistently exhibit lower performance, even for high-capacity models such as GPT-5-mini and Gemini-2.5-Flash. This discrepancy highlights a gap between superficial pattern recognition and deeper mechanistic understanding. We also observe that risk assessment scores are intermediate: models can often identify interactions, but fail to accurately assess their severity. The radar plot further shows that compositional reasoning benefits from increased context, yet performance on mechanistic and evidence-alignment tasks remains bottlenecked. Collectively, these results suggest that current LLMs possess robust outcome-prediction capabilities but are limited in faithful mechanistic and evidence-based reasoning. Overall, this capability breakdown underscores the necessity of multi-dimensional benchmarks like SafeDrug, which not only evaluate predictive accuracy but also probe reasoning fidelity, evidence use, and compositional generalization under complex pharmacological scenarios.




Figure 3: Capability breakdown of representative LLMs on SafeDrug.




Figure 4: Performance comparison (Set-F1) of evaluated LLMs on population-specific ADE prediction (Child, Adult, Older Adult).

3.4 Compositional Reasoning under Polypharmacy

We evaluate LLM performance on multi-drug prediction and risk assessment tasks to probe compositional reasoning under polypharmacy, where interactions among multiple drugs can give rise to complex adverse events. Exact Match scores are reported for 3, 4, 5, and 10 drug combinations, as summarized in Table 3 and visualized in Figure 7 (Appendix B.1.3). Across models, we observe that prediction accuracy generally increases or remains stable as the number of drugs grows, indicating that models can capture surface-level associations between drugs and potential interactions. For example, Gemini-2.5-Flash and GPT-5-mini achieve over 65% Exact Match on 10-drug prediction, demonstrating strong outcome-level pattern recognition. However, risk assessment performance is more variable and often lower than prediction, highlighting that models struggle to correctly estimate interaction severity as combinatorial complexity increases. Some models over-predict high-risk labels when the number of drugs is large, suggesting reliance on heuristics rather than explicit reasoning over specific interactions. These results indicate that while current LLMs can handle outcome-level multi-drug predictions, they remain limited in compositional reasoning, particularly in assessing risk. The findings underscore the need for benchmarks like SafeDrug that explicitly probe multi-drug interactions and compositional reasoning, beyond simple predictive associations.

3.5 Population-Specific ADE Evaluation

To assess model robustness across different patient populations, we evaluate LLM performance on ADE prediction for children (< 18 years old), adults (18–64 years old), and older adults (≥ 65 years old), as shown in Figure 4. This analysis is critical because pharmacological responses can vary with age, and overrepresentation of adult data in existing sources may bias model




Figure 5: Evaluation of LLM performance on evidence tasks in SafeDrug.

predictions. The results reveal a consistent adult-centric bias across most models. Models generally achieve higher metric scores on adult instances, with GPT-5-mini and Gemini-2.5-Flash reaching the highest performance. In contrast, performance on pediatric and older adult populations is substantially lower, particularly for smaller open-source models such as Qwen2.5-7B and LLaMA3.2-3B. This indicates that models tend to rely on patterns learned from adult-centric data and fail to generalize effectively to underrepresented populations. These findings highlight the importance of including population-stratified evaluation in benchmarks for safety-critical pharmacological reasoning. Without

such assessment, models may appear performant overall while exhibiting systematic biases that could compromise reliability in clinical settings for children and older adults.

3.6 Evidence Grounding and Faithfulness

We evaluate LLM performance on evidence-grounded tasks in SafeDrug, focusing on two dimensions: *evidence selection*, measured by Set-F1, and *evidence generation*, assessed via LLM-as-a-judge pass rate (Figure 5). Results indicate that models achieve moderate performance on evidence selection, with Set-F1 scores ranging from 0.22 (Gemini-2.5-Flash) to 0.55 (GPT-5-mini), suggesting that models often fail to reliably retrieve all relevant supporting information. Evidence generation performance is higher, with LLM-as-a-judge pass rates exceeding 0.9 for most models, indicating that when generating textual explanations, models can produce mostly plausible and coherent evidence. These findings reveal a gap between retrieval and generative reasoning: LLMs can produce convincing explanations, but they frequently omit or misalign supporting evidence. This highlights the need for benchmarks like SafeDrug that probe not only predictive accuracy but also the faithfulness of mechanistic and evidence-grounded reasoning.

3.7 Error Taxonomy Analysis

We analyze LLM failures across SafeDrug task categories using a taxonomy of five error types: context misalignment, missed signals (false negatives), hallucination/fabrication, over-risk/exaggeration, and superficial reasoning (Figure 6). The analysis reveals distinct failure patterns depending on the task. Prediction and risk assessment errors are dominated by missed signals and over-risk estimates, whereas mechanistic and evidence-grounded reasoning tasks show higher rates of context misalignment and hallucinated mechanisms. Compositional reasoning errors combine hallucination and missed signal effects, highlighting the challenges of multi-drug interactions. These findings demonstrate that SafeDrug effectively exposes systematic limitations in LLM reasoning, distinguishing superficial pattern recognition from mechanistic and evidence-grounded reasoning failures.




Figure 6: Errors are categorized into context misalignment, missed signals, hallucination/fabrication, over-risk/exaggeration, and superficial reasoning. Prediction and risk assessment errors are dominated by missed signals and over-risk estimates, whereas mechanistic and evidence-grounded reasoning tasks show higher rates of context misalignment and hallucinated mechanisms.

4 Discussion

Our experiments on SafeDrug show that LLMs perform well on outcome-level prediction but struggle with mechanistic reasoning, evidence grounding, risk assessment, and population-sensitive tasks, revealing a gap between surface associations and evidence-based reasoning for clinical safety. Future work should improve mechanistic and compositional reasoning, enhance evidence use, and address population biases via better data and modeling strategies. Current limitations include underrepresented subpopulations and incomplete evidence annotations, motivating expanded coverage and richer mechanistic data. Overall, SafeDrug provides a structured benchmark to advance robust and interpretable pharmacological reasoning in LLMs.

References

- [1] Karan Singhal, Shekoofeh Azizi, Tao Tu, S Sara Mahdavi, Jason Wei, Hyung Won Chung, Nathan Scales, Ajay Tanwani, Heather Cole-Lewis, Stephen Pfohl, et al. Large language models encode clinical knowledge. *Nature*, 620(7972):172–180, 2023.
- [2] Juexiao Zhou, Haoyang Li, Siyuan Chen, Zhangtianyi Chen, Zhongyi Han, and Xin Gao. Large language models in biomedicine and healthcare. *npj Artificial Intelligence*, 1(1):44, 2025.
- [3] Karan Singhal, Tao Tu, Juraj Gottweis, Rory Sayres, Ellery Wulczyn, Mohamed Amin, Le Hou, Kevin Clark, Stephen R Pfohl, Heather Cole-Lewis, et al. Toward expert-level medical question answering with large language models. *Nature medicine*, 31(3):943–950, 2025.
- [4] Shuyue S Li, Vidhisha Balachandran, Shangbin Feng, Jonathan S Ilgen, Emma Pierson, Pang W Koh, and Yulia Tsvetkov. Mediq: Question-asking llms and a benchmark for reliable interactive clinical reasoning. *Advances in Neural Information Processing Systems*, 37:28858–28888, 2024.
- [5] Akari Asai, Jacqueline He, Rulin Shao, Weijia Shi, Amanpreet Singh, Joseph Chee Chang, Kyle Lo, Luca Soldaini, Sergey Feldman, Mike D’Arcy, et al. Synthesizing scientific literature with retrieval-augmented language models. *Nature*, pages 1–7, 2026.
- [6] Shijie Wang, Wenqi Fan, Yue Feng, Lin Shanru, Xinyu Ma, Shuaiqiang Wang, and Dawei Yin. Knowledge graph retrieval-augmented generation for llm-based recommendation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics*, pages 27152–27168, 2025.
- [7] Anthony Yazdani, Alban Borner, Philipp Khlebnikov, Boya Zhang, Hossein Rouhizadeh, Poorya Amini, and Douglas Teodoro. An evaluation benchmark for adverse drug event prediction from clinical trial results. *Scientific Data*, 12(1):424, 2025.
- [8] Tengfei Ma, Kun Chen, Yongsheng Zang, Yujie Chen, Xuanbai Ren, Bosheng Song, Yiping Liu, et al. Self-supervised blending structural context of visual molecules for robust drug interaction prediction. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025.
- [9] Yuqin He, Tengfei Ma, Chaoyi Li, Pengsen Ma, Hongxin Xiang, Jianmin Wang, Yiping Liu, Bosheng Song, and Xiangxiang Zeng. Imageddi: Image-enhanced molecular motif sequence representation for drug-drug interaction prediction. *Information Fusion*, page 103574, 2025.
- [10] Craig Knox, Mike Wilson, Christen M Klinger, Mark Franklin, Eponine Oler, Alex Wilson, Allison Pon, Jordan Cox, Na Eun Chin, Seth A Strawbridge, et al. Drugbank 6.0: the drugbank knowledgebase for 2024. *Nucleic Acids Research*, 52(D1):D1265–D1275, 2024.
- [11] Marinka Zitnik, Monica Agrawal, and Jure Leskovec. Modeling polypharmacy side effects with graph convolutional networks. *Bioinformatics*, 34(13):i457–i466, 2018.
- [12] Joe G Greener and David T Jones. Differentiable molecular simulation can learn all the parameters in a coarse-grained force field for proteins. *PLoS One*, 16(9):e0256990, 2021.
- [13] Zhenqian Shen, Mingyang Zhou, Yongqi Zhang, and Quanming Yao. Benchmarking drug–drug interaction prediction methods: a perspective of distribution changes. *Bioinformatics*, 41(11):btaf569, 2025.
- [14] Xinmo Jin, Bowen Fan, Xunkai Li, Henan Sun, YuXin Zeng, Zekai Chen, Yuxuan Sun, Jia Li, Qiangqiang Dai, Hongchao Qin, et al. Openddi: A comprehensive benchmark for ddi prediction. *arXiv preprint arXiv:2602.00539*, 2026.
- [15] Yao Tian, Jiakai Yi, Ningning Wang, Chengkun Wu, Jinfu Peng, Shao Liu, Guoping Yang, and Dongsheng Cao. Ddinter 2.0: an enhanced drug interaction resource with expanded data coverage, new interaction types, and improved user interface. *Nucleic Acids Research*, 53(D1):D1356–D1362, 2025.

- [16] Bor Kavčič, Gašper Tkačik, and Tobias Bollenbach. Mechanisms of drug interactions between translation-inhibiting antibiotics. *Nature Communications*, 11(1):4013, 2020.
- [17] Junmo Kim, Joo Seong Kim, Ji-Hyang Lee, Min-Gyu Kim, Taehyun Kim, Chaeun Cho, Rae Woong Park, and Kwangsoo Kim. Pretrained patient trajectories for adverse drug event prediction using common data model-based electronic health records. *Communications Medicine*, 5(1):232, 2025.
- [18] Yaqing Wang, Zaifei Yang, and Quanming Yao. Accurate and interpretable drug-drug interaction prediction enabled by knowledge subgraph learning. *Communications Medicine*, 4(1):59, 2024.
- [19] Zhonghao Ren, Xiangxiang Zeng, Yizhen Lao, Zhuhong You, Yifan Shang, Quan Zou, and Chen Lin. Predicting rare drug-drug interaction events with dual-granular structure-adaptive and pair variational representation. *Nature Communications*, 16(1):3997, 2025.
- [20] Di Jin, Eileen Pan, Nassim Oufattole, Wei-Hung Weng, Hanyi Fang, and Peter Szolovits. What disease does this patient have? a large-scale open domain question answering dataset from medical exams. *Applied Sciences*, 11(14):6421, 2021.
- [21] Qiao Jin, Bhuwan Dhingra, Zhengping Liu, William Cohen, and Xinghua Lu. Pubmedqa: A dataset for biomedical research question answering. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2567–2577, 2019.
- [22] Anastasia Krithara, Anastasios Nentidis, Konstantinos Bougiatiotis, and Georgios Paliouras. Bioasq-qa: A manually curated corpus for biomedical question answering. *Scientific Data*, 10(1):170, 2023.
- [23] Michael Kuhn, Ivica Letunic, Lars Juhl Jensen, and Peer Bork. The sider database of drugs and side effects. *Nucleic Acids Research*, 44(D1):D1075–D1079, 2016.
- [24] Nicholas P Tatonetti, Patrick P Ye, Roxana Daneshjou, and Russ B Altman. Data-driven prediction of drug effects and interactions. *Science Translational Medicine*, 4(125):125ra31–125ra31, 2012.
- [25] Yao Tian, Jiakai Yi, Kun Li, Jinfu Peng, Youchao Deng, Dejun Jiang, and Dongsheng Cao. Machine learning prediction of pediatric adverse drug reactions using consensus-derived scarce data. *Communications Chemistry*, 2025.
- [26] Zhaoying Wang, Yingdan Shi, Xiang Liu, Can Chen, Jun Wen, and Ren Wang. Hoddi: A dataset of high-order drug-drug interactions for computational pharmacovigilance. *arXiv preprint arXiv:2502.06274*, 2025.
- [27] Aixin Liu, Aoxue Mei, Bangcai Lin, Bing Xue, Bingxuan Wang, Bingzheng Xu, Bochao Wu, Bowei Zhang, Chaofan Lin, Chen Dong, et al. Deepseek-v3. 2: Pushing the frontier of open large language models. *arXiv preprint arXiv:2512.02556*, 2025.
- [28] Binyuan Hui, Jian Yang, Zeyu Cui, Jiayi Yang, Dayiheng Liu, Lei Zhang, Tianyu Liu, Jiajun Zhang, Bowen Yu, Keming Lu, et al. Qwen2. 5-coder technical report. *arXiv preprint arXiv:2409.12186*, 2024.
- [29] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- [30] Aixin Liu, Bei Feng, Bing Xue, Bingxuan Wang, Bochao Wu, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, et al. Deepseek-v3 technical report. *arXiv preprint arXiv:2412.19437*, 2024.
- [31] Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card. *arXiv preprint arXiv:2410.21276*, 2024.

- [32] Gheorghe Comanici, Eric Bieber, Mike Schaeckermann, Ice Pasupat, Noveen Sachdeva, Inderjit Dhillon, Marcel Blistein, Ori Ram, Dan Zhang, Evan Rosen, et al. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv preprint arXiv:2507.06261*, 2025.
- [33] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- [34] Aaditya Singh, Adam Fry, Adam Perelman, Adam Tart, Adi Ganesh, Ahmed El-Kishky, Aidan McLaughlin, Aiden Low, AJ Ostrow, Akhila Ananthram, et al. Openai gpt-5 system card. *arXiv preprint arXiv:2601.03267*, 2025.
- [35] World Health Organization et al. Anatomical therapeutic chemical (atc) classification system, 2018.
- [36] Jae Yong Ryu, Hyun Uk Kim, and Sang Yup Lee. Deep learning improves prediction of drug–drug and drug–food interactions. *Proceedings of the National Academy of Sciences*, 115(18):E4304–E4311, 2018.
- [37] Xu Chu, Yang Lin, Yasha Wang, Leye Wang, Jiangtao Wang, and Jingyue Gao. Mlrda: A multi-task semi-supervised learning framework for drug-drug interaction prediction. In *Proceedings of the 28th International Joint Conference on Artificial Intelligence*, pages 4518–4524, 2019.
- [38] Jinhyuk Lee, Wonjin Yoon, Sungdong Kim, Donghyeon Kim, Sunkyu Kim, Chan Ho So, and Jaewoo Kang. Biobert: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 36(4):1234–1240, 2020.
- [39] Yue Yu, Kexin Huang, Chao Zhang, Lucas M Glass, Jimeng Sun, and Cao Xiao. Sumgnn: multi-typed drug interaction prediction via efficient knowledge graph summarization. *Bioinformatics*, 37(18):2988–2995, 2021.
- [40] Xuan Lin, Zhe Quan, Zhi-Jie Wang, Tengfei Ma, and Xiangxiang Zeng. Kggnn: Knowledge graph neural network for drug-drug interaction prediction. In *Proceedings of the 29th International Joint Conference on Artificial Intelligence*, volume 380, pages 2739–2745, 2020.
- [41] Yujie Chen, Tengfei Ma, Xixi Yang, Jianmin Wang, Bosheng Song, and Xiangxiang Zeng. Muffin: multi-scale feature fusion for drug–drug interaction prediction. *Bioinformatics*, 37(17):2651–2658, 2021.
- [42] Yongqi Zhang, Quanming Yao, Ling Yue, Xian Wu, Ziheng Zhang, Zhenxi Lin, and Yefeng Zheng. Emerging drug interaction prediction enabled by a flow-based graph neural network with biomedical network. *Nature Computational Science*, 3(12):1023–1033, 2023.
- [43] Haotong Du, Quanming Yao, Juzheng Zhang, Yang Liu, and Zhen Wang. Customized subgraph selection and encoding for drug–drug interaction prediction. *Advances in Neural Information Processing Systems*, 37:109582–109608, 2024.

Technical Appendix

A Details of Data Preprocessing

A.1 Mechanism and Evidence Supervision

Mechanism labels and evidence rationales were not collected through large-scale manual annotation from scratch. Instead, we first derived supervision from DDInter 2.0 itself, including interaction descriptions, management recommendations, reference lists, and alternative-drug fields. After evidence-aware sampling, all selected source interactions were manually checked by our team for pair consistency, task assignment, label validity, and evidence quality. Thus, the benchmark is best characterized as *source-grounded with manual verification*, rather than fully human-authored. At the benchmark level, all non-generative supervision is source-derived rather than LLM-generated.

This includes classification, multi-label, multi-choice, and span-selection targets. In our cleaned release, this accounts for 2,154 of 2,554 instances ($\approx 84.3\%$). LLM assistance is used only for generative reference fields, namely mechanism reasoning and evidence reasoning, which together account for 400 of 2,554 instances ($\approx 15.7\%$). These fields are generated from DDInter text and, when available, literature-enriched abstracts, and are treated as reference answers rather than hidden labels for structured prediction. To guard against circularity, we adopt three safeguards. First, LLM-synthesized fields are never used as gold labels for objective or semi-structured tasks. Second, the released `model_input` hides answer-bearing fields for objective tasks, reducing shortcut access to supervision. Third, generative evaluation is reported separately, and LLM-based judging is performed with a fixed external judge rather than using the structured labels themselves. As a result, LLM assistance serves only as a controlled augmentation for free-text reference construction, not as a source of circular supervision for the main structured benchmark.

A.2 Therapy-level Drug Replacement and Clinical Plausibility

For therapy-level drug replacement (or substitution) tasks, SafeDrug follows a structured and clinically motivated construction strategy based on DDInter2.0. Specifically, replacement candidates are derived using the Anatomical Therapeutic Chemical (ATC) classification system [35], where drugs sharing the same **fourth-level ATC code** are considered therapeutically equivalent. This level of the ATC hierarchy corresponds to pharmacological or therapeutic subgroups, ensuring that candidate replacements have similar mechanisms of action and clinical indications.

This design enables scalable and consistent annotation of replacement options while maintaining a reasonable degree of clinical plausibility. Compared to heuristic or model-inferred substitutions, ATC-based grouping provides a standardized and widely accepted framework for identifying alternative drugs within the same therapeutic class. For SafeDrug-small, all replacement candidates are further manually verified by domain experts. The verification process ensures that suggested alternatives are clinically plausible, do not introduce obvious safety risks, and are consistent with the context of the interaction. Ambiguous or inappropriate substitutions are removed during this process.

We emphasize that these replacement annotations are intended for *benchmarking and evaluation purposes only*. They approximate therapy-level alternatives based on pharmacological similarity, but do not constitute clinical treatment recommendations. Real-world drug substitution requires comprehensive consideration of patient-specific factors, comorbidities, and clinical guidelines.

A.3 Details of SafeDrug-large

The **SafeDrug-large** benchmark is systematically designed to evaluate Large Language Models (LLMs) on diverse and complex pharmacovigilance tasks. In this section, we detail the demographic categorization principles and provide concrete task formulations through QA examples.

A.3.1 Population Demographics and Data Sources

A key differentiator of SafeDrug-large is its fine-grained demographic stratification, explicitly integrating patient age into Adverse Drug Event (ADE) predictions.

To construct the age-specific subset, we primarily utilize data from the **CT-ADE** and **Pediatric** databases. We standardize the population into three distinct age categories: *child*, *adult*, and *older_adult*. The categorization follows these strict rules:

- **CT-ADE Mapping:** The inherent age attributes provided in the CT-ADE database naturally align with our three categories (*child*, *adult*, *older_adult*), which we directly adopt.
- **Pediatric Database Mapping:** All records sourced from the specialized Pediatric dataset are systematically mapped and merged into the *child* category.

To preserve this rich metadata, our global integrated database (`Single_Drug_Global_Data.jsonl`) organizes demographic attributes under a comprehensive `demographics` key. This structured schema contains several sub-attributes, including `age_group`, `gender`, `is_healthy_volunteer`, and `is_pediatric_specific`. This granular structural design ensures that downstream tasks accurately reflect demographic-aware pharmacological reasoning.

565 A.3.2 Task Formulation and QA Examples

566 To clearly illustrate the evaluation formats, we provide representative JSON-formatted QA examples
567 for each fundamental task in the SafeDrug-large benchmark. All prompts explicitly instruct the LLMs
568 to strictly follow output constraints to facilitate automated evaluation.

569 1. Age-Specific ADE Prediction

570 Evaluates ADE identification constrained by a specific age demographic (e.g., Adult).

Example from ade_prediction_age/adult.jsonl

```
571 {  
572 1 "task_name": "Adult_SideEffect_DBQA",  
573 2 "prompt": "You are an expert clinical pharmacologist... \n[SCENARIO:  
574 3 Adult Case Study]\nPatient Profile: Adult Patient (18-64 years)  
575 .\nDrug Prescribed: Ulevostinag\nTask: Identify side effects  
576 that are specifically documented or observed in adult  
577 populations for this drug.\nOptions:\nA. Infusion site  
578 thrombosis\nB. Injection site reaction\n...\nF. None of the  
579 above\nInstructions: Provide ONLY the option letters.",  
580 "answer": "B, E"  
581 4 }  
582 5 }  
583
```

584 2. ADE Risk and Frequency Prediction

585 Predicts the occurrence frequency (CIOMS definitions) of an ADE based on SIDER and OFFSIDES.

Example from ade_risk_prediction/sider/sider_risk_dataset.jsonl

```
586 {  
587 1 "task_name": "SIDER_Frequency_Classification",  
588 2 "prompt": "Drug Prescribed: montelukast... \nAdverse Drug Event (ADE  
589 3 ): Hypersensitivity\nTask: Predict the occurrence frequency  
590 category of this specific ADE.\nOptions:\nA. Very common (>=  
591 10%)\nB. Common (1% to < 10%)\nC. Uncommon (0.1% to < 1%)\nD.  
592 Rare (0.01% to < 0.1%)\nE. Very rare (< 0.01%)",  
593 "answer": "D"  
594 4 }  
595 5 }  
596
```

597 3. Dual-Drug Interaction (DDI) Type Prediction

598 Identifies the primary pharmacological mechanism or type of interaction between two drugs.

Example from ddi_prediction/ddi_prediction.jsonl

```
599 {  
600 1 "task_name": "DDI_Type_Prediction",  
601 2 "prompt": "[SCENARIO: Drug-Drug Interaction Assessment]\nDrug A:  
602 3 Methylergometrine\nDrug B: Ciprofloxacin\nTask: Identify the  
603 primary pharmacological interaction type or mechanism.\nOptions  
604 :\nA. Others\nB. Synergy\nC. Unknown\nD. Excretion\nE.  
605 Metabolism",  
606 "answer": "E"  
607 4 }  
608 5 }  
609
```

610 4. Dual-Drug Interaction Risk Assessment

611 Predicts the clinical severity level of a specific dual-drug interaction.

Example from ddi_risk_prediction/ddi_risk_prediction.jsonl

```
612 {  
613 1 "task_name": "DDI_Severity_Prediction",  
614 2 "prompt": "Drug A: Succinylcholine\nDrug B: Nitrous oxide\  
615 3 nInteraction Type: Synergy\nTask: Predict the clinical severity  
616 (risk level) of this specific type of interaction.\nOptions:\nA.  
617 Major\nB. Minor\nC. Moderate",  
618
```

```

619 4   "answer": "C"
620 5 }

```

622 5. Polypharmacy (Multi-Drug) Interaction

623 Determines whether concurrent administration of multiple drugs (e.g., 3, 4, 5, or 10) causes significant
624 ADEs.

Example from multiple_drug_prediction/ddi_prediction/3_drug.jsonl

```

625 {
626 1 {
627 2   "task_name": "Polypharmacy_Interaction_Prediction_3_Drugs",
628 3   "prompt": "Drugs Prescribed Concurrently:\n1. Dasabuvir\n2. Isopropyl
629             alcohol\n3. Ribociclib\nTask: Predict whether the concurrent
630             administration of all these specific drugs together leads to a
631             known significant adverse drug-drug interaction (DDI).\nOptions
632             :\nYes\nNo",
633 4   "answer": "No"
634 5 }
635 }

```

636 A.4 Details of SafeDrug-small

637 Our benchmark is constructed from DDInter 2.0 through a multi-stage pipeline. We start with 277K
638 interaction records containing structured fields such as descriptions, management recommendations,
639 references, and alternative drugs, and select 200 evidence-rich interactions via stratified sampling
640 based on an evidence-oriented score incorporating reference density, severity, management complete-
641 ness, and mechanistic cues. All sampled instances are manually verified to ensure consistency of
642 interaction pairs, labels, and evidence, and ambiguous or noisy samples are removed. From each
643 verified interaction, we construct a structured evaluation set spanning three task tiers: objective
644 (e.g., interaction check, severity classification), semi-structured (e.g., evidence and mechanism span
645 selection), and generative (e.g., mechanism and evidence reasoning), with task-dependent context to
646 avoid label leakage.

647 To further ensure reliability, we construct SafeDrug-small via stratified sampling of 3,000 instances
648 from SafeDrug-large based on evidence density, followed by expert verification by a team consisting of
649 a clinical practitioner (Master-level), a Ph.D. researcher in AI for health, and a professor specializing
650 in AI for biological systems. The verification follows concise guidelines:

651 [Guidelines]

- 652 (1) Interaction correctness: check drug pair, type, and severity consistency.
- 653 (2) Evidence alignment: ensure references support the interaction.
- 654 (3) Mechanism validity: confirm biological plausibility.
- 655 (4) Quality control: remove ambiguous, inconsistent, or low-confidence cases.

656 This process yields 2,554 high-confidence instances. We further augment `interaction_check`
657 with synthetic hard negatives to enable true binary evaluation. The final benchmark contains 2,554
658 instances, split by source interaction to prevent data leakage (1,817/331/406 for train/dev/test).
659 The released `split-clean` version retains only evaluation-critical fields, including `model_input`,
660 `label`, `answer`, and `scoring_hint`, making it directly usable for controlled evaluation.

661 B Additional Experiments

662 B.1 Evaluation Protocol

663 In this section, we provide the formal mathematical definitions for the three metric families utilized
664 in our evaluation framework.

665 B.1.1 Classification Metrics (Prediction-oriented Tasks)

666 Let $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$ be the dataset with N samples, where $y_i \in \mathcal{C}$ is the ground truth label and $\hat{y}_i$
667 is the predicted label. For a specific class c , we denote True Positives (TP_c), False Positives (FP_c),
668 True Negatives (TN_c), and False Negatives (FN_c).

Table 4: Sample counts for each task in SafeDrug-large.

Task	Samples
ADE Prediction (Age)	20,040
ADE Prediction (Overall)	38,934
ADE Risk Prediction (OFFSIDES)	67,452
ADE Risk Prediction (SIDER)	70,290
DDI Prediction	260,060
DDI Risk Prediction	260,060
DDI Reasoning: Counterfactual	1,076
DDI Reasoning: Mechanism Generation	260,060
DDI Reasoning: Mechanism Prediction	260,060
DDI Reasoning: Substitution Prediction	415,413
Multiple-Drug DDI Prediction (3 drugs)	84,376
Multiple-Drug DDI Prediction (4 drugs)	35,116
Multiple-Drug DDI Prediction (5 drugs)	11,508
Multiple-Drug DDI Prediction (10 drugs)	601
Multiple-Drug Risk Prediction (3 drugs)	113,595
Multiple-Drug Risk Prediction (4 drugs)	34,413
Multiple-Drug Risk Prediction (5 drugs)	10,474
Multiple-Drug Risk Prediction (10 drugs)	417
Total	1,943,945

669 **Accuracy and F1-Scores.** Accuracy measures the overall proportion of correct predictions:

$$\text{Accuracy} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{y}_i = y_i) \quad (1)$$

670 where $\mathbb{I}(\cdot)$ is the indicator function. For a specific class c (or binary classification), Precision (P_c)
671 and Recall (R_c) are defined as $P_c = \frac{TP_c}{TP_c + FP_c}$ and $R_c = \frac{TP_c}{TP_c + FN_c}$. The **Binary-F1** score is the
672 harmonic mean of precision and recall for the positive class:

$$\text{F1}_c = 2 \cdot \frac{P_c \cdot R_c}{P_c + R_c} \quad (2)$$

673 To address class imbalance, we employ **Macro-F1** (unweighted mean of F1 across all $|\mathcal{C}|$ classes)
674 and **Micro-F1** (F1 computed over globally aggregated TP , FP , and FN):

$$\text{Macro-F1} = \frac{1}{|\mathcal{C}|} \sum_{c=1}^{|\mathcal{C}|} \text{F1}_c, \quad \text{Micro-F1} = \frac{2 \sum_c TP_c}{2 \sum_c TP_c + \sum_c FP_c + \sum_c FN_c} \quad (3)$$

675 **ROC-AUC.** The Area Under the Receiver Operating Characteristic Curve (ROC-AUC) evaluates
676 the model’s discriminative ability across all decision thresholds. Mathematically, it is the probability
677 that a randomly chosen positive sample x^+ is assigned a higher predicted probability $\hat{p}$ than a
678 randomly chosen negative sample x^- :

$$\text{ROC-AUC} = \mathbb{E}_{x^+, x^-} [\mathbb{I}(\hat{p}(x^+) > \hat{p}(x^-))] \quad (4)$$

679 B.1.2 Set-based Metrics (Multiple Valid Outputs)

680 For tasks yielding a set of elements (e.g., adverse drug event sets), let Y denote the ground truth set
681 and $\hat{Y}$ denote the predicted set.

682 **Exact Match (EM) and Partial Match.** **Exact Match** requires the predicted set to be strictly
683 identical to the target set:

$$\text{EM} = \mathbb{I}(\hat{Y} = Y) \quad (5)$$

684 **Partial Match** is quantified using the Jaccard Similarity, which measures the degree of overlap:

$$\text{Partial Match (Jaccard)} = \frac{|\hat{Y} \cap Y|}{|\hat{Y} \cup Y|} \quad (6)$$




Figure 7: Multiple-drug prediction performance under increasing polypharmacy complexity. Exact Match scores indicate that models can capture outcome-level interactions, but their performance does not necessarily reflect mechanistic or evidence-based reasoning.

Set-F1. Set-F1 adapts the standard F1-score for unranked sets. We define Set-Precision $P_{set} = \frac{|\hat{Y} \cap Y|}{|\hat{Y}|}$ and Set-Recall $R_{set} = \frac{|\hat{Y} \cap Y|}{|Y|}$. The Set-F1 is formulated as:

$$\text{Set-F1} = 2 \cdot \frac{P_{set} \cdot R_{set}}{P_{set} + R_{set}} = \frac{2|\hat{Y} \cap Y|}{|\hat{Y}| + |Y|} \quad (7)$$

B.1.3 Generation Metrics (Open-ended Reasoning)

For text generation tasks, let $\mathbf{r} = (r_1, \dots, r_L)$ be the reference sequence and $\mathbf{c} = (c_1, \dots, c_M)$ be the generated candidate sequence.

BLEU and ROUGE. BLEU evaluates precision-based n -gram overlap penalized by a brevity penalty (BP):

$$\text{BLEU} = \text{BP} \cdot \exp \left(\sum_{n=1}^N w_n \log p_n \right), \quad \text{where } \text{BP} = \begin{cases} 1 & \text{if } M > L \\ e^{1-L/M} & \text{if } M \leq L \end{cases} \quad (8)$$

where p_n is the modified n -gram precision. Conversely, **ROUGE** (specifically ROUGE-L) focuses on recall via the Longest Common Subsequence (LCS):

$$R_{lcs} = \frac{|\text{LCS}(\mathbf{r}, \mathbf{c})|}{L}, \quad P_{lcs} = \frac{|\text{LCS}(\mathbf{r}, \mathbf{c})|}{M}, \quad \text{ROUGE-L} = \frac{(1 + \beta^2) R_{lcs} P_{lcs}}{R_{lcs} + \beta^2 P_{lcs}} \quad (9)$$

LLM-as-a-judge. To capture semantic consistency beyond lexical overlap, we define an evaluation function $\mathcal{J}_\theta$ parameterized by a strong language model. Given an input query q , the reference answer $\mathbf{r}$, and the candidate answer $\mathbf{c}$, the judge yields a binary or scalar score s :

$$s = \mathcal{J}_\theta(\text{prompt_template}(q, \mathbf{r}, \mathbf{c})) \quad (10)$$

where $s \in [0, 1]$ represents semantic alignment, factual correctness, and reasoning coherence.

To improve evaluation consistency, the judge operates under structured prompts with task-specific checklists and constrained outputs. Detailed calibration strategies and reliability considerations are provided in Appendix C.

B.2 Analysis of Polypharmacy Prediction Trends

As shown in Figure 7, most models exhibit stable or improving performance as the number of drugs increases, suggesting that LLMs can capture aggregated interaction patterns at the outcome




Figure 8: Fine-grained evaluation of LLMs on SafeDrug-small across objective, semi-structured, and generative tasks. While models perform well on interaction checking and generation-based reasoning, performance drops significantly on evidence selection and mechanism span alignment, highlighting a gap between fluent generation and evidence-grounded reasoning.

level. However, this trend does not necessarily reflect true compositional reasoning, as higher scores may arise from memorized co-occurrence patterns rather than mechanistic understanding. Notably, several models show performance saturation or fluctuation under higher drug counts, indicating limited robustness in handling complex multi-drug interactions. These results further highlight the gap between predictive performance and reliable reasoning in polypharmacy settings.

B.3 Fine-grained Task Analysis on SafeDrug-small

We present a fine-grained evaluation across all task categories in SafeDrug-small, including objective, semi-structured, and generative tasks. Figure 8 reports model performance on representative tasks such as interaction check, severity classification, interaction type prediction, risk label selection, management action prediction, alternative drug selection, evidence selection, mechanism span selection, and reasoning-based tasks.

Objective Tasks. For objective tasks such as interaction check and alternative drug selection, most models achieve relatively higher performance, indicating that LLMs can capture explicit interaction patterns and structured decision rules. However, performance on severity classification and interaction type prediction remains moderate, suggesting limitations in distinguishing finer-grained clinical attributes.

Semi-structured Tasks. Performance drops noticeably on semi-structured tasks, including evidence selection, mechanism span selection, and management rationale selection. In particular, mechanism

span selection and rationale prediction show substantial degradation, reflecting the difficulty of aligning outputs with precise spans or structured evidence. This suggests that LLMs struggle to ground their predictions in specific supporting information, even when relevant context is provided.

Generative Reasoning Tasks. In contrast, models achieve high scores on mechanism reasoning and evidence reasoning when evaluated with LLM-as-a-judge. This indicates that models can produce fluent and seemingly plausible explanations. However, combined with the weaker performance on evidence selection and span alignment, this suggests that such explanations may not be faithfully grounded in the provided evidence, highlighting a gap between generation quality and factual grounding.

Key Observations. Overall, the results reveal a consistent pattern: LLMs perform well on surface-level prediction and free-form generation, but struggle with tasks requiring precise evidence alignment and structured reasoning. This discrepancy indicates that current models may rely on memorized patterns or heuristic reasoning rather than truly grounded, mechanism-aware inference.

C Evaluation Protocol Clarifications

To improve transparency and reproducibility, we provide additional details on key evaluation design choices, including error annotation, LLM-based judgment, and metric definitions.

C.1 Error Taxonomy and Annotation Strategy

Sampling Strategy. To conduct a robust qualitative error analysis, we designed a stratified random sampling strategy. We first aggregated all incorrect predictions (where Exact Match = 0) from the evaluated models across all datasets. These errors were then mapped into five macro-task families: *Prediction*, *Risk Assessment*, *Evidence Grounding*, *Mechanistic Reasoning*, and *Compositional Reasoning*. Using a fixed random seed to ensure reproducibility, we sampled up to **80 error instances per task family**, yielding a total error analysis pool of approximately 400 samples.

Annotation Process. To scale error attribution across this diverse pool, we adopted an **LLM-as-an-annotator** approach. The annotator is prompted under a strict expert clinical pharmacologist persona and provided with the question, ground truth, model prediction, and definitions of five mutually exclusive error categories (e.g., *Hallucination*, *Missed Signal*, *Over-Risk*, *Context Misalignment*, *Superficial Reasoning*).

Reliability Considerations. We ensure consistency through: (1) deterministic prompting with fixed templates, (2) mutually exclusive category design, (3) context-grounded justification requirements, and (4) fixed random seeds for reproducibility. This enables stable and scalable error attribution for macro-level analysis.

C.2 Calibration of LLM-as-a-Judge

For generative tasks, we employ an LLM-based judge to evaluate free-text outputs.

Calibration. The judge is constrained using **task-specific checklists** (e.g., causal explanation, evidence usage, clinical relevance) and grounded on both the reference answer and the source context. It outputs a strict binary score (1: correct/aligned, 0: incorrect or hallucinated), converting evaluation into deterministic checklist verification.

Reliability Considerations. We improve robustness via: (1) checklist-constrained evaluation, (2) explicit context grounding, (3) binary decision space, and (4) standardized prompt templates. These constraints reduce arbitrariness and improve reproducibility.

763 D Data Licenses and Redistribution Permissions

764 To answer the concerns regarding data redistribution, we confirm that all original data sources
765 incorporated in our study—including CT-ADE², DDInter2.0³, HODDI⁴, SIDER⁵, OFFSIDES⁶,
766 DrugBank⁷, and the pediatric resources⁸—are publicly accessible. Their respective licensing terms
767 explicitly permit usage for academic research and non-commercial purposes.

768 In strict adherence to the redistribution requirements of these foundational resources, we guarantee
769 that all derived instances, instruction-tuning datasets, and generated evidence constructed during
770 this study will be fully open-sourced. They will be officially released to the community under a
771 clear and standard non-commercial license (e.g., Creative Commons Attribution-NonCommercial
772 4.0 International, CC BY-NC 4.0). This ensures that our derived resources remain freely available to
773 support future academic and non-commercial research endeavors in the medical domain.

774 E Related Work

775 **Drug–Drug Interaction and Drug Safety Modeling.** Drug–drug interaction (DDI) prediction
776 has been extensively studied using machine learning and graph-based methods, aiming to identify
777 potential adverse interactions or beneficial drug combinations. Existing approaches typically model
778 DDI data as graphs and leverage drug features [36, 37, 38], knowledge graph embeddings [39, 40, 41],
779 or graph neural networks [11, 42, 43] to predict interaction types. While these methods have achieved
780 strong predictive performance, their evaluation is largely limited to classification settings, focusing on
781 accuracy rather than a deeper understanding of pharmacological mechanisms or safety implications.

782 **Benchmarking and Evaluation for DDI Prediction.** Recent work has recognized the limitations
783 of current evaluation protocols and proposed more systematic benchmarking efforts [7, 15]. For
784 example, prior studies highlight the lack of unified evaluation frameworks, inconsistencies in datasets
785 and metrics, and insufficient consideration of real-world scenarios such as interactions involving
786 new drugs [13, 14]. Furthermore, performance can vary significantly across different DDI types and
787 data distributions, and the role of side information (e.g., biomedical networks) remains insufficiently
788 understood. These findings suggest that existing benchmarks primarily capture predictive capability
789 but fall short in assessing robustness, generalization, and practical applicability.

790 **Limitations of Existing Benchmarks for Pharmacological Reasoning.** Despite progress in both
791 DDI prediction and LLM evaluation, there remains a fundamental gap between current benchmarks
792 and the requirements of real-world drug safety assessment. Existing DDI benchmarks emphasize
793 prediction accuracy but do not evaluate reasoning, evidence grounding, or clinical decision sup-
794 port. Conversely, LLM benchmarks focus on general QA and lack domain-specific structure for
795 pharmacological reasoning. Critically, current evaluation frameworks do not account for multi-drug
796 interactions, population heterogeneity, or the hierarchical nature of drug safety reasoning. As a
797 result, the ability of LLMs to support safety-critical pharmacological reasoning remains largely
798 underexplored.

799 F Broader Impact

800 SafeDrug aims to advance research on reliable and interpretable pharmacological reasoning in large
801 language models. By providing a large-scale benchmark with evidence grounding, mechanistic
802 annotations, and population-aware evaluation, our dataset can support the development of safer AI
803 systems for biomedical and clinical applications. We hope SafeDrug will facilitate research on robust
804 reasoning, trustworthy evidence utilization, and clinically informed decision support, particularly in
805 safety-critical settings such as adverse drug event prediction and drug–drug interaction analysis.

²<https://huggingface.co/anthonyyazdani1>

³<https://ddinter2.scbdd.com/>

⁴<https://github.com/TIML-Group/HODDI>

⁵<http://sideeffects.embl.de>

⁶<https://nsides.io/#offsides-and-twosides>

⁷<https://go.drugbank.com/>

⁸<https://github.com/TY-CSU/Predicting-Adverse-Drug-Reactions-in-Children>

806 At the same time, our work also carries potential risks. Models evaluated or trained on SafeDrug may
807 generate plausible but incorrect pharmacological recommendations, especially in complex clinical
808 scenarios involving polypharmacy, rare populations, or incomplete evidence. Such outputs could be
809 harmful if used directly in real-world clinical decision-making without expert supervision. In addition,
810 biases inherited from existing biomedical databases and literature may lead to uneven performance
811 across demographic groups or disease conditions.

812 To mitigate these risks, SafeDrug is released strictly for research and benchmarking purposes rather
813 than clinical deployment. We explicitly emphasize that benchmark outputs should not be interpreted
814 as medical advice or treatment recommendations. Furthermore, we include evidence-grounded tasks
815 and population-aware evaluation to encourage future research on transparency, fairness, and reliability
816 in biomedical LLM systems.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction clearly state the contributions of SafeDrug, including dataset construction, task design, and evaluation findings, which are supported by experimental results in Sections 3.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: We discuss limitations in Section 4, including underrepresented populations, incomplete evidence annotations, and scope constraints of pharmacological datasets.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: This work is empirical and does not include theoretical results.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: We provide detailed descriptions of dataset construction, task formulation, evaluation metrics, and experimental settings in Sections 3, Appendix, and the anonymous data warehouse, enabling reproducibility.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: We provide the SafeDrug dataset and evaluation scripts, along with instructions for data preparation and benchmarking, in the anonymous data warehouse.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.

974	• At submission time, to preserve anonymity, the authors should release anonymized
975	versions (if applicable).
976	• Providing as much information as possible in supplemental material (appended to the
977	paper) is recommended, but including URLs to data and code is permitted.
978	6. Experimental setting/details
979	Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
980	rameters, how they were chosen, type of optimizer) necessary to understand the results?
981	Answer: [Yes]
982	Justification: We specify model settings, zero-shot evaluation protocol, data splits, and
983	preprocessing details in Section 3 and the Appendix.
984	Guidelines:
985	• The answer [N/A] means that the paper does not include experiments.
986	• The experimental setting should be presented in the core of the paper to a level of detail
987	that is necessary to appreciate the results and make sense of them.
988	• The full details can be provided either with the code, in appendix, or as supplemental
989	material.
990	7. Experiment statistical significance
991	Question: Does the paper report error bars suitably and correctly defined or other appropriate
992	information about the statistical significance of the experiments?
993	Answer: [No]
994	Justification: We do not report error bars as experiments are deterministic zero-shot evalua-
995	tions without stochastic training; variability is limited and does not affect conclusions.
996	Guidelines:
997	• The answer [N/A] means that the paper does not include experiments.
998	• The authors should answer [Yes] if the results are accompanied by error bars, confidence
999	intervals, or statistical significance tests, at least for the experiments that support the
1000	main claims of the paper.
1001	• The factors of variability that the error bars are capturing should be clearly stated (for
1002	example, train/test split, initialization, random drawing of some parameter, or overall
1003	run with given experimental conditions).
1004	• The method for calculating the error bars should be explained (closed form formula,
1005	call to a library function, bootstrap, etc.)
1006	• The assumptions made should be given (e.g., Normally distributed errors).
1007	• It should be clear whether the error bar is the standard deviation or the standard error
1008	of the mean.
1009	• It is OK to report 1-sigma error bars, but one should state it. The authors should
1010	preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis
1011	of Normality of errors is not verified.
1012	• For asymmetric distributions, the authors should be careful not to show in tables or
1013	figures symmetric error bars that would yield results that are out of range (e.g., negative
1014	error rates).
1015	• If error bars are reported in tables or plots, the authors should explain in the text how
1016	they were calculated and reference the corresponding figures or tables in the text.
1017	8. Experiments compute resources
1018	Question: For each experiment, does the paper provide sufficient information on the com-
1019	puter resources (type of compute workers, memory, time of execution) needed to reproduce
1020	the experiments?
1021	Answer: [Yes]
1022	Justification: We describe compute resources, including GPU type and API-based model
1023	access, in Section 3 and the Appendix.
1024	Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research adheres to the NeurIPS Code of Ethics, and all data sources are publicly available or curated without involving personal data.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: We discuss both positive impacts (improving drug safety evaluation) and potential risks (misuse of model predictions) in the Appendix F.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [\[Yes\]](#)

Justification: The dataset is intended for research use only, and we explicitly discourage direct clinical deployment without expert supervision.

Guidelines:

- The answer [\[N/A\]](#) means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: All datasets used (e.g., DrugBank, SIDER, DDInter) are properly cited, and their usage follows respective licenses and terms.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: SafeDrug is a new dataset, and we provide detailed documentation on its construction, format, and usage in the paper and appendix.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.

- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: This work does not involve crowdsourcing or human subject experiments.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: No human subjects are involved, and thus IRB approval is not required.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: LLMs are used as core components for evaluation and for processing literature-derived evidence (e.g., DeepSeek), as described in Sections 3.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.